

# Theodore Zhang

AI/ML Engineer | 408-680-1290 | Cupertino, California | zhangtheodore1234@gmail.com

## EDUCATION

### Arizona State University

B.S. Computer Science

Aug 2022 – May 2026

Tempe, AZ

## EXPERIENCE

### HandShake

AI Evaluation Engineer

April 2026 – June 2026

Tempe, AZ

- Designed adversarial test suites targeting model reasoning and code generation, surfacing systematic failure modes across edge-case inputs.
- Developed "Golden Solutions" for complex engineering problems, establishing Large Language Model (LLM) performance benchmarks to measure generation quality and consistency.

Technologies used: Python, LLM Evaluation, Adversarial Testing

### ASML

DevOps Engineer Intern

May 2025 – August 2025

San Jose, CA

- Built CI/CD pipelines automating AI model integration into production workflows, reducing manual handoff steps and accelerating release cycles.
- Provisioned infrastructure-as-code using Azure Bicep, standardizing cloud environments and cutting provisioning time across AI research and production teams.
- Configured network authentication securing cloud-hosted ML resources; partnered with AI research teams to align DevOps practices with model deployment requirements.

Technologies used: Azure Bicep, Microsoft Azure, Azure DevOps, CI/CD, YAML, Linux

### ASML

Software Quality Assurance Engineer Intern

May 2024 – August 2024

San Jose, CA

- Audited 1,000+ Python files for accuracy, functionality, and security compliance, establishing quality standards for production scripts.
- Implemented Machine Learning (ML) models to analyze Confluence test and design documents, automatically flagging missing information and critical gaps for engineering teams.
- Conducted automated and manual testing of Python scripts, validating functionality, usability, and security across the Software Development Life Cycle (SDLC).

Technologies used: Python, Llama 3.1, Machine Learning, Git

## PROJECTS

### KimuntuPro AI

Aug 2025 – May 2026

- Designed a Retrieval-Augmented Generation (RAG) legal assistant spanning 9 practice areas, integrating a custom categorized knowledge base with GPT-4.1-mini for structured JSON output and fallback handling.
- Integrated real-time web search to supplement RAG retrieval, enabling the assistant to surface up-to-date legal information and answer questions beyond the static knowledge base.
- Engineered a full-stack Next.js application with a Role-Based Access Control (RBAC) system, mapping subscription tiers to feature flags and protected routes.
- Built a real-time admin dashboard using Firebase and Recharts to visualize user metrics, AI usage, and revenue.

### AI-Powered Document Intelligence System

May 2025 – August 2025

- Designed a RAG-based system ingesting PDF, HTML, DOCX, PPTX, XLSX, images, TXT, JSON, and CSV.
- Implemented inline citation generation for complete source transparency.

### Test & Design Plan Analyzer (Machine Learning)

May 2024 – August 2024

- Built a Python analyzer using fine-tuned Llama 3.1 to evaluate Confluence documents, identifying missing test cases and surfacing critical gaps for engineering review.

## TECHNICAL SKILLS

Languages: Python, C++, Java, Lua, SQL, TypeScript, English, Mandarin

Technologies: Git, Docker, Azure DevOps, Microsoft Azure, Linux, Firebase, Next.js, LangChain, YAML, Bicep, LLM Evaluation & Benchmarking, RAG Pipelines, Fine-Tuning (Llama 3.1), Adversarial Testing, Web Search Integration